

---

# CSE 151B Project Milestone Report

---

Alexander Friend  
apfriend@ucsd.edu

## 1 Task Description and Exploratory Analysis

### 1.1 Problem A [1 points]

The task for this Kaggle competition is to predict the positions of 60 individual vehicles 3 seconds into the future, given an initial 2 second observation. This is an important task as autonomous vehicles (AV) are being increasingly rolled out to the public, and are expected to become the future standard of automobile transportation. In order for this future to be realized however, AV's must be able to predict future movement and positions of objects in their vicinity with high accuracy in order for AV's to be safer than human drivers.

Our data is split into two sets, a training set and test set, titled `new_train` and `new_val_in`, respectively. The training dataset contains the following fields:

- `p_in` - the (x,y) position input in the first two seconds (19 time steps)
- `v_in` - the (x,y) velocity input in the first two seconds (19 time steps)
- `p_out` - the (x,y) position output in the next three seconds (30 time steps)
- `v_out` - the (x,y) velocity output in the next three seconds (30 time steps)
- `track_id` - the `track_id` for each vehicle in the output sequence (30 time steps)
- `scene_idx` - the id of this scene
- `agent_id` - track id for the agent in this scene
- `car_mask` - boolean index for the real car, we need to align the car numbers
- `lane` - (x,y,z) for centerline nodes in this scene (z position is always 0)
- `lane_norm` - (x,y,z) the direction of each lane node (z position is always 0)

The test set contains the same fields, except it is lacking the `p_out` and `v_out` fields, as these are the fields to be predicted.

Using these datasets, our input will be the the positions (`p_in`), velocities (`v_in`), and ID's of 60 other vehicles (`track_id`) in the scene, as well as information on the position of the lanes. Lane information will include the position of the center of the lanes (`lane`) in the scene and the direction of each lane (`lane_norm`). There can be more than one lane in a scene.

The output will be the ID of each of the 60 other vehicles in the scene (`scene_idx`) with their respective positions (`p_out`). The output will have 10 records/second for the 3 second prediction period for one of the 60 vehicles for each test set. This means that the output csv file will have  $(10 \times 3 \times 60) + 1 = 3201$ , rows including the header, as each row contains the `scene_idx`, and `p_out` of all the vehicles in that time frame. It will therefore have 61 columns, one column for each vehicle position (`p_out`) as well as a column for the corresponding scene id (`scene_idx`).

## 1.2 Problem B [1 points]

The training data set contains 205,944 training values, each of dimension  $2 \times 60 \times 19 \times 4$ . Each training value contains the input and output positions and velocities of each of the 60 cars 1.9 seconds, at a sampling rate of 10Hz.

The validation set contains 3,200 values, each of dimension  $60 \times 19 \times 4$ . Each validation value contains only input positions and velocities of each of the 60 cars for 1.9 seconds, at a sampling rate of 10Hz.

Statistic	$x$ -Position	$y$ -Position	$x$ -Velocity	$y$ -Velocity
count	2.280000e+06	2.280000e+06	2.280000e+06	2.280000e+06
mean	2.108387e+02	3.245734e+02	3.604328e-02	-1.889331e-02
std	7.032012e+02	8.485306e+02	1.749104e+00	2.185653e+00
min	0.000000e+00	0.000000e+00	-1.144689e+02	-1.688118e+02
25	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
50	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
75	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
max	4.735943e+03	4.092156e+03	7.909621e+01	1.461265e+02

Table 1: Table of input positions and velocities of sample

Statistic	$x$ -Position	$y$ -Position	$x$ -Velocity	$y$ -Velocity
count	3.600000e+06	3.600000e+06	3.600000e+06	3.600000e+06
mean	2.109253e+02	3.245277e+02	3.481067e-02	-1.777745e-02
std	7.034579e+02	8.483126e+02	1.797118e+00	2.223628e+00
min	0.000000e+00	0.000000e+00	1.008539e+02	-1.215008e+02
25	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
50	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
75	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
max	4.736574e+03	4.092117e+03	8.774132e+01	1.484342e+02

Table 2: Table of Output positions and velocities of sample




Figure 1: Histogram of the X and Y input positions from a subsample of 2000 training files




Figure 2: Histogram of the X and Y output positions from a subsample of 2000 training files

Analyzing these inputs we find that the vast majority of  $(x, y)$  input and output positions are at position 0, as seen in Figure 1, Figure 2, and Table 1 above. Interestingly, the  $x$ -positions seem to be more closely clustered around 0, while the  $y$ -positions have more variance, meaning that most cars are moving mostly in the  $y$ -dimension, with less overall travel in the  $x$ -dimension. This trend carries through the beginning and end of the scene (the input and output data).

Looking at only the magnitudes of velocities shows the same trend, with the vast majority of velocities being 0, however there is also slightly more variance in the distribution of velocities in the  $y$ -dimension compared to the  $x$ -dimension.

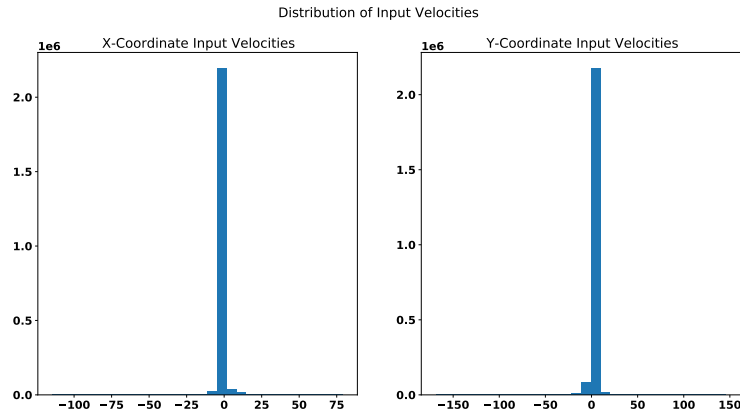


Figure 3: Histogram of the X and Y input velocities from a subsample of 2000 training files

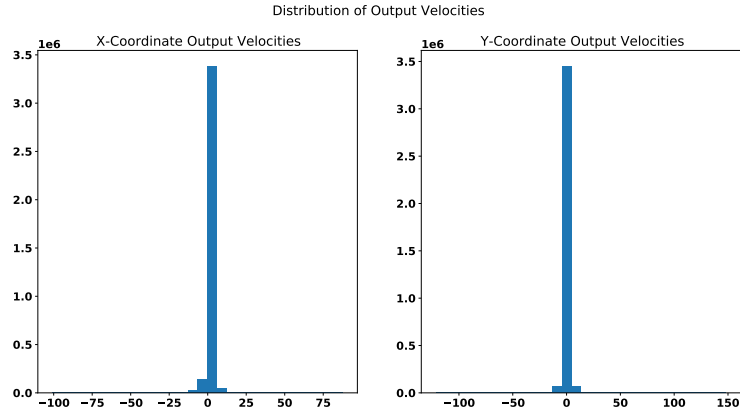


Figure 4: Histogram of the X and Y output velocities from a subsample of 2000 training files

Taking a look at the velocity distributions in Figure 3, Figure 4), and Table 2 we again see that the velocities are centered around 0 in both the  $x$  and  $y$  dimensions at both the beginning and end of each scene. Unlike the position distribution however, velocities can be negative, which is why the distribution is balanced on both sides.

## 2 Deep Learning Model and Experiment Design

### 2.1 Problem A [1 Points]

Thus far my best performing model is a PyTorch linear neural net running in an Anaconda environment. The platform I am currently working in is my local machine running Ubuntu 20.04 with a 4-core 4-thread Intel i7-7600k CPU running at 4.2 GHz, a GTX 1070 GPU with a max clock speed of 1721 MHz and 8 GB of GDDR5 memory, and 16 GB of 3000 MHz DDR4 memory.

For my initial baseline model I made a simple single layer linear neural network, using Adam as an optimizer and mean square error as my loss function. I used a learning rate of 0.001, and trained my best model for 10 epochs. Each epoch took approximately 3 minutes to train.

I chose mean square error as it increases the penalty for worse predictions, compared to a optimizer like mean absolute error, which penalizes with a linear rate.

In order to get the correct outputs from my trained model, I sliced the output prediction tensors so as to only keep the first inner 2 elements of each tensor (the predicted  $x$  and  $y$  coordinates), and then selected the correct car based on each inputs `agent_id`.

### 2.2 Problem B [1 Points]

I experimented with multilayer linear models and an LSTM model as well, however these actually performed worse on the validation set once I uploaded them to Kaggle. Thus far, the best performing model is a linear one, which takes in the assumption that cars will be moving at about the same rate over the 5 seconds period of training and prediction.

### 3 Experiment Results and Future Work

#### 3.1 Problem A [1 points]




Figure 5: Plot of the average training mean square error for my best performing single layer model

This model however achieved a significantly lower loss on the training set compared to the final validation set on Kaggle, suggesting that my model has overfit the training data. This means that a more complex model with more variance will need to be implemented.

Based on my lack of success with an LSTM approach, I will continue to work on more complex linear models, as well as trying other architectures such as a RNN model, which may achieve a better result as it can take into account previous training examples.

My current rank on the Kaggle leaderboard is 40 out of 57.

#### A Appendix

<https://github.com/apfriend/cse151b-kaggle.git>